

## ✓ SETUP

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount

```
%cd /content/drive/MyDrive/Capstone/Data_Cleaning
```

```
/content/drive/MyDrive/Capstone/Data_Cleaning
```

```
%pwd
```

```
' /content/drive/MyDrive/Capstone/Data_Cleaning'
```

```
%ls
```

|                                        |                                           |
|----------------------------------------|-------------------------------------------|
| 'Copy of Analysis&ModellingTest.ipynb' | DeploymentTest.ipynb                      |
| 'Copy of DataCleaning.ipynb'           | <a href="#">gradio_crash_app_package/</a> |
| crash_file.csv                         | sms_input.csv                             |
| DataCleaning.ipynb                     | UPS_Crash_Risk_Modeling.ipynb             |

## ✓ DATA CLEANING

```
import pandas as pd
import numpy as np
```

```
df = pd.read_csv("crash_file.csv")
df1 = pd.read_csv("sms_input.csv")
```

df.head()

|   | CHANGE_DATE      | CRASH_ID | REPORT_STATE | REPORT_NUMBER | REPORT_DATE | REPC |
|---|------------------|----------|--------------|---------------|-------------|------|
| 0 | 19950927<br>0000 | 837419   | NH           | NH0003808-93  | 19930301    |      |
| 1 | 19960212<br>0000 | 676123   | MA           | MA0000000138  | 19940315    |      |
| 2 | 19960212<br>0000 | 676732   | MA           | MA0000000798  | 19951024    |      |
| 3 | 19960329<br>0000 | 676874   | MA           | MA0000000940  | 19950706    |      |
| 4 | 19970208<br>0000 | 677185   | MA           | MA0000001251  | 19960917    |      |

5 rows x 59 columns

df.columns

```
Index(['CHANGE_DATE', 'CRASH_ID', 'REPORT_STATE', 'REPORT_NUMBER',
      'REPORT_DATE', 'REPORT_TIME', 'REPORT_SEQ_NO', 'DOT_NUMBER',
      'CI_STATUS_CODE', 'FINAL_STATUS_DATE', 'LOCATION', 'CITY_CODE',
      'CITY',
      'STATE', 'COUNTY_CODE', 'TRUCK_BUS_IND', 'TRAFFICWAY_ID',
      'ACCESS_CONTROL_ID', 'ROAD_SURFACE_CONDITION_ID',
      'CARGO_BODY_TYPE_ID',
      'GVW_RATING_ID', 'VEHICLE_IDENTIFICATION_NUMBER',
      'VEHICLE_LICENSE_NUMBER', 'VEHICLE_LIC_STATE',
      'VEHICLE_HAZMAT_PLACARD',
      'WEATHER_CONDITION_ID', 'VEHICLE_CONFIGURATION_ID',
      'LIGHT_CONDITION_ID', 'HAZMAT_RELEASED', 'AGENCY',
      'VEHICLES_IN_ACCIDENT', 'FATALITIES', 'INJURIES', 'TOW_AWAY',
      'FEDERAL_RECORDABLE', 'STATE_RECORDABLE',
      'SNET_VERSION_NUMBER',
      'SNET_SEQUENCE_ID', 'TRANSACTION_CODE', 'TRANSACTION_DATE',
      'UPLOAD_FIRST_BYTE', 'UPLOAD_DOT_NUMBER',
      'UPLOAD_SEARCH_INDICATOR',
      'UPLOAD_DATE', 'ADD_DATE', 'CRASH_CARRIER_ID',
      'CRASH_CARRIER_NAME',
      'CRASH_CARRIER_STREET', 'CRASH_CARRIER_CITY',
      'CRASH_CARRIER_CITY_CODE',
      'CRASH_CARRIER_STATE', 'CRASH_CARRIER_ZIP_CODE',
      'CRASH_COLONIA',
      'DOCKET_NUMBER', 'CRASH_CARRIER_INTERSTATE', 'NO_ID_FLAG',
      'STATE_NUMBER', 'STATE_ISSUING_NUMBER',
      'CRASH_EVENT_SEQ_ID_DESC'],
      dtype='object')
```

df1.head()

|   | Report_number | Report_seq_no | DOT_Number | Report_Date | Report_State |
|---|---------------|---------------|------------|-------------|--------------|
| 0 | AL0004636248  | 1             | 21800      | 04-APR-24   | AL           |
| 1 | AL0004640085  | 1             | 21800      | 17-APR-24   | AL           |
| 2 | AL0004645287  | 1             | 21800      | 30-APR-24   | AL           |
| 3 | AL0004646793  | 2             | 21800      | 04-MAY-24   | AL           |
| 4 | AL0004661083  | 1             | 21800      | 10-JUN-24   | AL           |

5 rows x 22 columns

df1.columns

Index(['Report\_number', 'Report\_seq\_no', 'DOT\_Number', 'Report\_Date', 'Report\_State', 'Fatalities', 'Injuries', 'Tow\_Away', 'Hazmat\_Released', 'Trafficway\_Desc', 'Access\_Control\_Desc', 'Road\_Surface\_Condition\_Desc', 'Weather\_Condition\_Desc', 'Light\_Condition\_Desc', 'Vehicle\_ID\_Number', 'Vehicle\_License\_number', 'Vehicle\_License\_State', 'Severity\_Weight', 'Time\_Weight', 'Citation\_Issued\_Desc', 'Seq\_Num', 'Not\_Preventable'], dtype='object')

df.info()

|   |               |               |         |
|---|---------------|---------------|---------|
| 2 | REPORT_STATE  | 6525 non-null | object  |
| 3 | REPORT_NUMBER | 6525 non-null | object  |
| 4 | REPORT_DATE   | 6525 non-null | int64   |
| 5 | REPORT_TIME   | 6525 non-null | int64   |
| 6 | REPORT_SEQ_NO | 6525 non-null | int64   |
| 7 | DOT_NUMBER    | 239 non-null  | float64 |

|    |                               |      |          |         |
|----|-------------------------------|------|----------|---------|
| 8  | CI_STATUS_CODE                | 6525 | non-null | object  |
| 9  | FINAL_STATUS_DATE             | 553  | non-null | float64 |
| 10 | LOCATION                      | 6496 | non-null | object  |
| 11 | CITY_CODE                     | 3060 | non-null | float64 |
| 12 | CITY                          | 6521 | non-null | object  |
| 13 | STATE                         | 6525 | non-null | object  |
| 14 | COUNTY_CODE                   | 6525 | non-null | int64   |
| 15 | TRUCK_BUS_IND                 | 6525 | non-null | object  |
| 16 | TRAFFICWAY_ID                 | 6165 | non-null | float64 |
| 17 | ACCESS_CONTROL_ID             | 236  | non-null | float64 |
| 18 | ROAD_SURFACE_CONDITION_ID     | 6397 | non-null | float64 |
| 19 | CARGO_BODY_TYPE_ID            | 5844 | non-null | float64 |
| 20 | GVW_RATING_ID                 | 5982 | non-null | float64 |
| 21 | VEHICLE_IDENTIFICATION_NUMBER | 5621 | non-null | object  |
| 22 | VEHICLE_LICENSE_NUMBER        | 6510 | non-null | object  |
| 23 | VEHICLE_LIC_STATE             | 6523 | non-null | object  |
| 24 | VEHICLE_HAZMAT_PLACARD        | 1637 | non-null | object  |
| 25 | WEATHER_CONDITION_ID          | 6383 | non-null | float64 |
| 26 | VEHICLE_CONFIGURATION_ID      | 6429 | non-null | float64 |
| 27 | LIGHT_CONDITION_ID            | 6406 | non-null | float64 |
| 28 | HAZMAT_RELEASED               | 3540 | non-null | object  |
| 29 | AGENCY                        | 6437 | non-null | object  |
| 30 | VEHICLES_IN_ACCIDENT          | 6448 | non-null | float64 |
| 31 | FATALITIES                    | 6525 | non-null | int64   |
| 32 | INJURIES                      | 6525 | non-null | int64   |
| 33 | TOW_AWAY                      | 6525 | non-null | object  |
| 34 | FEDERAL_RECORDABLE            | 6525 | non-null | object  |
| 35 | STATE_RECORDABLE              | 6525 | non-null | object  |
| 36 | SNET_VERSION_NUMBER           | 6180 | non-null | object  |
| 37 | SNET_SEQUENCE_ID              | 6525 | non-null | int64   |
| 38 | TRANSACTION_CODE              | 6525 | non-null | object  |
| 39 | TRANSACTION_DATE              | 6525 | non-null | object  |
| 40 | UPLOAD_FIRST_BYTE             | 6057 | non-null | object  |
| 41 | UPLOAD_DOT_NUMBER             | 182  | non-null | object  |
| 42 | UPLOAD_SEARCH_INDICATOR       | 5965 | non-null | object  |
| 43 | UPLOAD_DATE                   | 6525 | non-null | object  |
| 44 | ADD_DATE                      | 6525 | non-null | object  |
| 45 | CRASH_CARRIER_ID              | 6525 | non-null | int64   |
| 46 | CRASH_CARRIER_NAME            | 6525 | non-null | object  |
| 47 | CRASH_CARRIER_STREET          | 6525 | non-null | object  |
| 48 | CRASH_CARRIER_CITY            | 6525 | non-null | object  |
| 49 | CRASH_CARRIER_CITY_CODE       | 0    | non-null | float64 |
| 50 | CRASH_CARRIER_STATE           | 6525 | non-null | object  |
| 51 | CRASH_CARRIER_ZIP_CODE        | 6525 | non-null | object  |
| 52 | CRASH_COLONIA                 | 0    | non-null | float64 |
| 53 | DOCKET_NUMBER                 | 99   | non-null | float64 |
| 54 | CRASH_CARRIER_INTERSTATE      | 6525 | non-null | object  |
| 55 | NO_ID_FLAG                    | 6314 | non-null | float64 |
| 56 | STATE_NUMBER                  | 6495 | non-null | object  |
| 57 | STATE_ISSUING_NUMBER          | 6525 | non-null | object  |

```
58 CRASH_EVENT_SEQ_ID_DESC          6471 non-null    object
dtypes: float64(16), int64(9), object(34)
memory usage: 2.9+ MB
```

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2311 entries, 0 to 2310
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Report_number                        2311 non-null   object
1   Report_seq_no                       2311 non-null   int64
2   DOT_Number                         2311 non-null   int64
3   Report_Date                         2311 non-null   object
4   Report_State                       2311 non-null   object
5   Fatalities                         2311 non-null   int64
6   Injuries                          2311 non-null   int64
7   Tow_Away                          2311 non-null   bool
8   Hazmat_Released                   1488 non-null   object
9   Trafficway_Desc                   2123 non-null   object
10  Access_Control_Desc               1436 non-null   object
11  Road_Surface_Condition_Desc       2303 non-null   object
12  Weather_Condition_Desc           2294 non-null   object
13  Light_Condition_Desc             2310 non-null   object
14  Vehicle_ID_Number                2300 non-null   object
15  Vehicle_License_number           2301 non-null   object
16  Vehicle_License_State            2305 non-null   object
17  Severity_Weight                  2303 non-null   float64
18  Time_Weight                     2303 non-null   float64
19  Citation_Issued_Desc             2287 non-null   object
20  Seq_Num                         2311 non-null   int64
21  Not_Preventable                  1 non-null     object
dtypes: bool(1), float64(2), int64(5), object(14)
memory usage: 381.5+ KB
```

## ✓ Standardize Column Names

```
def clean_columns(df):
    df.columns = (
        df.columns
        .str.lower()
        .str.strip()
        .str.replace(" ", "_")
        .str.replace("-", "_")
    )
    return df

df = clean_columns(df)
df1 = clean_columns(df1)
```

```
df.head(5)
```

|   | change_date      | crash_id | report_state | report_number | report_date | repo |
|---|------------------|----------|--------------|---------------|-------------|------|
| 0 | 19950927<br>0000 | 837419   | NH           | NH0003808-93  | 19930301    |      |
| 1 | 19960212<br>0000 | 676123   | MA           | MA0000000138  | 19940315    |      |
| 2 | 19960212<br>0000 | 676732   | MA           | MA0000000798  | 19951024    |      |
| 3 | 19960329<br>0000 | 676874   | MA           | MA0000000940  | 19950706    |      |
| 4 | 19970208<br>0000 | 677185   | MA           | MA0000001251  | 19960917    |      |

5 rows x 59 columns

```
df1.head(3)
```

|   | report_number | report_seq_no | dot_number | report_date | report_state |
|---|---------------|---------------|------------|-------------|--------------|
| 0 | AL0004636248  | 1             | 21800      | 04-APR-24   | AL           |
| 1 | AL0004640085  | 1             | 21800      | 17-APR-24   | AL           |
| 2 | AL0004645287  | 1             | 21800      | 30-APR-24   | AL           |

3 rows x 22 columns



df.columns

```
Index(['change_date', 'crash_id', 'report_state', 'report_number',
      'report_date', 'report_time', 'report_seq_no', 'dot_number',
      'ci_status_code', 'final_status_date', 'location', 'city_code',
      'city',
      'state', 'county_code', 'truck_bus_ind', 'trafficway_id',
      'access_control_id', 'road_surface_condition_id',
      'cargo_body_type_id',
      'gvw_rating_id', 'vehicle_identification_number',
      'vehicle_license_number', 'vehicle_lic_state',
      'vehicle_hazmat_placard',
      'weather_condition_id', 'vehicle_configuration_id',
      'light_condition_id', 'hazmat_released', 'agency',
      'vehicles_in_accident', 'fatalities', 'injuries', 'tow_away',
      'federal_recordable', 'state_recordable',
      'snet_version_number',
      'snet_sequence_id', 'transaction_code', 'transaction_date',
      'upload_first_byte', 'upload_dot_number',
      'upload_search_indicator',
      'upload_date', 'add_date', 'crash_carrier_id',
      'crash_carrier_name',
      'crash_carrier_street', 'crash_carrier_city',
      'crash_carrier_city_code',
      'crash_carrier_state', 'crash_carrier_zip_code',
      'crash_colonia',
      'docket_number', 'crash_carrier_interstate', 'no_id_flag',
      'state_number', 'state_issuing_number',
      'crash_event_seq_id_desc'],
      dtype='object')
```

df1.columns

```
Index(['report_number', 'report_seq_no', 'dot_number', 'report_date',
      'report_state', 'fatalities', 'injuries', 'tow_away',
      'hazmat_released',
      'trafficway_desc', 'access_control_desc',
      'road_surface_condition_desc',
      'weather_condition_desc', 'light_condition_desc',
      'vehicle_id_number',
      'vehicle_license_number', 'vehicle_license_state',
      'severity_weight',
      'time_weight', 'citation_issued_desc', 'seq_num',
      'not_preventable'],
      dtype='object')
```

```
print(df.shape)
print(df1.shape)
```

```
(6525, 59)
(2311, 22)
```

## ✓ Cleaning on df

```
df.info()
```

```

2    report_state      6525 non-null  object
3    report_number    6525 non-null  object
4    report_date      6525 non-null  int64
5    report_time      6525 non-null  int64
6    report_seq_no    6525 non-null  int64
7    dot_number       239 non-null   float64
8    ci_status_code   6525 non-null  object
9    final_status_date 553 non-null   float64
10   location        6496 non-null  object
11   city_code       3060 non-null  float64
12   city            6521 non-null  object
13   state           6525 non-null  object
14   county_code     6525 non-null  int64
15   truck_bus_ind   6525 non-null  object
16   trafficway_id   6165 non-null  float64
17   access_control_id 236 non-null   float64
18   road_surface_condition_id 6397 non-null float64
19   cargo_body_type_id 5844 non-null float64
20   gvw_rating_id   5982 non-null  float64
21   vehicle_identification_number 5621 non-null object
22   vehicle_license_number 6510 non-null object
23   vehicle_lic_state 6523 non-null  object
24   vehicle_hazmat_placard 1637 non-null object
25   weather_condition_id 6383 non-null float64
26   vehicle_configuration_id 6429 non-null float64
27   light_condition_id 6406 non-null float64
28   hazmat_released 3540 non-null  object
29   agency           6437 non-null  object
30   vehicles_in_accident 6448 non-null float64
31   fatalities       6525 non-null  int64
32   injuries         6525 non-null  int64
33   tow_away         6525 non-null  object
34   federal_recordable 6525 non-null  object
35   state_recordable 6525 non-null  object
36   snet_version_number 6180 non-null object
37   snet_sequence_id 6525 non-null  int64
```

```

38 transaction_code      6525 non-null object
39 transaction_date      6525 non-null object
40 upload_first_byte     6057 non-null object
41 upload_dot_number     182 non-null object
42 upload_search_indicator 5965 non-null object
43 upload_date           6525 non-null object
44 add_date              6525 non-null object
45 crash_carrier_id      6525 non-null int64
46 crash_carrier_name    6525 non-null object
47 crash_carrier_street  6525 non-null object
48 crash_carrier_city    6525 non-null object
49 crash_carrier_city_code 0 non-null float64
50 crash_carrier_state    6525 non-null object
51 crash_carrier_zip_code 6525 non-null object
52 crash_colonia         0 non-null float64
53 docket_number         99 non-null float64
54 crash_carrier_interstate 6525 non-null object
55 no_id_flag            6314 non-null float64
56 state_number          6495 non-null object
57 state_issuing_number  6525 non-null object
58 crash_event_seq_id_desc 6471 non-null object
dtypes: float64(16), int64(9), object(34)
memory usage: 2.9+ MB

```

```
df.head()
```

|   | change_date      | crash_id | report_state | report_number | report_date | repor |
|---|------------------|----------|--------------|---------------|-------------|-------|
| 0 | 19950927<br>0000 | 837419   | NH           | NH0003808-93  | 19930301    |       |
| 1 | 19960212<br>0000 | 676123   | MA           | MA0000000138  | 19940315    |       |
| 2 | 19960212<br>0000 | 676732   | MA           | MA0000000798  | 19951024    |       |
| 3 | 19960329<br>0000 | 676874   | MA           | MA0000000940  | 19950706    |       |
| 4 | 19970208<br>0000 | 677185   | MA           | MA0000001251  | 19960917    |       |

5 rows x 59 columns

## Handle Date/Time

```
# Preserve original (always good practice)
df['report_date_raw'] = df['report_date']
df['report_time_raw'] = df['report_time']

# --- Fix report_date ---
df['report_date'] = pd.to_datetime(
    df['report_date'].astype(str),
    format='%Y%m%d',
    errors='coerce'
)

# --- Fix report_time ---
# Ensure 4-digit format (e.g., 30 -> "0030")
df['report_time'] = df['report_time'].astype(str).str.zfill(4)

df['report_time'] = pd.to_datetime(
    df['report_time'],
    format='%H%M',
    errors='coerce'
).dt.time
```

```
df[['report_date_raw', 'report_date', 'report_time_raw', 'report_time
```

|   | report_date_raw | report_date | report_time_raw | report_time |  |
|---|-----------------|-------------|-----------------|-------------|-------------------------------------------------------------------------------------|
| 0 | 19930301        | 1993-03-01  | 1430            | 14:30:00    |                                                                                     |
| 1 | 19940315        | 1994-03-15  | 1550            | 15:50:00    |                                                                                     |
| 2 | 19951024        | 1995-10-24  | 1404            | 14:04:00    |                                                                                     |
| 3 | 19950706        | 1995-07-06  | 907             | 09:07:00    |                                                                                     |
| 4 | 19960917        | 1996-09-17  | 1622            | 16:22:00    |                                                                                     |
| 5 | 20010928        | 2001-09-28  | 1450            | 14:50:00    |                                                                                     |
| 6 | 20011120        | 2001-11-20  | 1000            | 10:00:00    |                                                                                     |
| 7 | 19981220        | 1998-12-20  | 447             | 04:47:00    |                                                                                     |
| 8 | 20020807        | 2002-08-07  | 1207            | 12:07:00    |                                                                                     |
| 9 | 20020828        | 2002-08-28  | 1245            | 12:45:00    |                                                                                     |

```
df['report_date'].isna().sum()
df['report_time'].isna().sum()
```

```
np.int64(1)
```

### ***Fix Remaining Structural Data Type Issues***

```
# --- Convert date columns ---
date_cols = [
    'change_date',
    'transaction_date',
    'upload_date',
    'add_date'
]

for col in date_cols:
    df[col] = pd.to_datetime(df[col], errors='coerce')

# --- Convert categorical ID columns ---
cat_cols = [
    'trafficway_id',
    'access_control_id',
    'road_surface_condition_id',
    'weather_condition_id',
    'vehicle_configuration_id',
    'light_condition_id',
    'cargo_body_type_id',
    'gvw_rating_id'
]

for col in cat_cols:
    df[col] = df[col].astype('category')
```

```
df[date_cols].dtypes
df[cat_cols].dtypes
```

```

0
trafficway_id      category
access_control_id  category
road_surface_condition_id  category
weather_condition_id  category
vehicle_configuration_id  category
light_condition_id  category
cargo_body_type_id  category
gvw_rating_id      category
```

```
dtype: object
```

```
missing_summary = pd.DataFrame({
    'column': df.columns,
    'missing_count': df.isna().sum(),
    'missing_pct': (df.isna().sum() / len(df)) * 100
}).sort_values(by='missing_pct', ascending=False)
```

```
missing_summary.head(20)
```

|                         | column                  | missing_count | missing_pct |
|-------------------------|-------------------------|---------------|-------------|
| crash_carrier_city_code | crash_carrier_city_code | 6525          | 100.0       |
| crash_colonia           | crash_colonia           | 6525          | 100.0       |
| docket_number           | docket_number           | 6426          | 98.4        |
| upload_dot_number       | upload_dot_number       | 6343          | 97.2        |
| access_control_id       | access_control_id       | 6289          | 96.3        |
| dot_number              | dot_number              | 6286          | 96.3        |
| final_status_date       | final_status_date       | 5972          | 91.5        |
| vehicle_hazmat_placard  | vehicle_hazmat_placard  | 4888          | 74.5        |
| city_code               | city_code               | 3405          | 50.8        |

|                               |                               |      |      |
|-------------------------------|-------------------------------|------|------|
| city_code                     | city_code                     | 3403 | 33.1 |
| hazmat_released               | hazmat_released               | 2985 | 45.7 |
| vehicle_identification_number | vehicle_identification_number | 904  | 13.8 |
| cargo_body_type_id            | cargo_body_type_id            | 681  | 10.4 |
| upload_search_indicator       | upload_search_indicator       | 560  | 8.5  |
| gvw_rating_id                 | gvw_rating_id                 | 543  | 8.3  |
| upload_first_byte             | upload_first_byte             | 468  | 7.1  |
| trafficway_id                 | trafficway_id                 | 360  | 5.5  |
| snet_version_number           | snet_version_number           | 345  | 5.2  |
| no_id_flag                    | no_id_flag                    | 211  | 3.2  |
| weather_condition_id          | weather_condition_id          | 142  | 2.1  |
| road_surface_condition_id     | road_surface_condition_id     | 128  | 1.9  |

Next steps:

[Generate code with missing\\_summary](#)

[New interactive sheet](#)

OBSERVATIONS

Completely Useless columns:

- crash\_carrier\_city\_code
- crash\_colonia

Best Clean Predictors:

- Weather
- Road surface
- Trafficway
- Vehicle type

| Type         | Action (later)        |
|--------------|-----------------------|
| 100% missing | Drop                  |
| >90% missing | Likely drop           |
| 40–75%       | Convert to "Unknown"  |
| <15%         | Keep + minor handling |



## ✓ df1 Handling

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2311 entries, 0 to 2310
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   report_number                        2311 non-null   object
1   report_seq_no                       2311 non-null   int64
2   dot_number                          2311 non-null   int64
3   report_date                         2311 non-null   object
4   report_state                        2311 non-null   object
5   fatalities                         2311 non-null   int64
6   injuries                          2311 non-null   int64
7   tow_away                          2311 non-null   bool
8   hazmat_released                    1488 non-null   object
9   trafficway_desc                    2123 non-null   object
10  access_control_desc                1436 non-null   object
11  road_surface_condition_desc        2303 non-null   object
12  weather_condition_desc             2294 non-null   object
13  light_condition_desc               2310 non-null   object
14  vehicle_id_number                  2300 non-null   object
15  vehicle_license_number             2301 non-null   object
16  vehicle_license_state              2305 non-null   object
17  severity_weight                    2303 non-null   float64
18  time_weight                        2303 non-null   float64
19  citation_issued_desc              2287 non-null   object
20  seq_num                           2311 non-null   int64
21  not_preventable                    1 non-null      object
dtypes: bool(1), float64(2), int64(5), object(14)
memory usage: 381.5+ KB
```

```
# --- Fix report_date ---
df1['report_date'] = pd.to_datetime(
    df1['report_date'],
    errors='coerce'
)

# --- Create severity_class ---
df1['severity_class'] = np.select(
    [
        df1['fatalities'] > 0,
        df1['injuries'] > 0
    ],
    [
        'Fatal',
        'Injury'
    ],
    default='Property_Damage'
)

df1['severity_class'] = df1['severity_class'].astype('category')

/tmp/ipykernel_15445/3669085816.py:2: UserWarning: Could not infer format
df1['report_date'] = pd.to_datetime(
```

```
df1[['fatalities', 'injuries', 'severity_class']].head()
df1['severity_class'].value_counts()
```

|                        | count |
|------------------------|-------|
| <b>severity_class</b>  |       |
| <b>Property_Damage</b> | 1425  |
| <b>Injury</b>          | 835   |
| <b>Fatal</b>           | 51    |

**dtype:** int64

```
cols_to_clean = [
    'weather_condition_desc',
    'light_condition_desc',
    'road_surface_condition_desc',
    'trafficway_desc',
    'hazmat_released'
]

for col in cols_to_clean:
    df1[col] = df1[col].astype(str).str.strip().str.lower()
```

```
cols_to_clean = [
    'weather_condition_desc',
    'light_condition_desc',
    'road_surface_condition_desc',
    'trafficway_desc',
    'hazmat_released'
]

for col in cols_to_clean:
    df1[col] = df1[col].astype(str).str.strip().str.lower()
```

```
for col in cols_to_clean:
    df1[col] = df1[col].replace('nan', 'unknown')
```

```
for col in cols_to_clean:
    print(f"\n{col}")
    print(df1[col].value_counts().head(10))
```

```
weather_condition_desc
weather_condition_desc
no adverse conditions      1791
rain                       219
other                      92
snow                       85
unknown                   69
fog                        21
blowing sand soil dirt or snow  14
severe crosswinds          10
sleet hail                 10
Name: count, dtype: int64
```

```
light_condition_desc
```

```

light_condition_desc
daylight                                1456
dark - not lighted                      411
dark - lighted                          341
dusk                                    47
dawn                                    38
dark - unknown roadway lighting         11
unknown                                7
Name: count, dtype: int64

road_surface_condition_desc
road_surface_condition_desc
dry                                    1806
wet                                    314
snow                                   79
ice                                    68
slush                                  16
unknown                                13
water(standing moving)                 11
other                                   3
sandmuddirtoil or gravel               1
Name: count, dtype: int64

trafficway_desc
trafficway_desc
two-way trafficway divided unprotected median  882
two-way trafficway not divided                 526
two-way trafficway divided positive barrier    520
unknown                                         263
one-way trafficway not divided                 110
not reported                                   10
Name: count, dtype: int64

hazmat_released
hazmat_released
false          1484
unknown        823
true           4
Name: count, dtype: int64

```

# Post Data CLeaning

◆ df1 (SMS Input) → Primary Dataset

Using this for:

- Severity (target)
- Visualizations
- Model later
- Web app inputs (VERY IMPORTANT)

◆ df (Crash File) → Secondary / Support Dataset

Use this for:

Extra exploration if needed Not required for main pipeline


## ✓ EXPLORATORY VISUALIZATIONS

*Distribution of crash severity*

```
import matplotlib.pyplot as plt

df1['severity_class'].value_counts().plot(
    kind='bar'
)

plt.title("Crash Severity Distribution")
plt.xlabel("Severity Class")
plt.ylabel("Count")
plt.xticks(rotation=0)
plt.show()
```



## Interpretation — Crash Severity Distribution

The distribution shows that the majority of crashes are Property Damage Only (PDO), followed by Injury crashes, while Fatal crashes are relatively rare.

Specifically:

Property Damage crashes dominate the dataset (60%+) Injury crashes form a substantial secondary category (35%) Fatal crashes represent a very small proportion (~2%) 🧠 What This Means (Important Insight) Crash severity is highly imbalanced

Most crashes are not severe Severe outcomes (fatal) are rare events Injury crashes are non-trivial There is a meaningful middle category This justifies using a multi-class model Fatal crashes are rare but critical Even though small in number, they are: High-impact Operationally important 🎯 Why This Matters for Your Project

This supports your framework because:

You are not predicting just “crash vs no crash” You are predicting severity levels, which is: More realistic More valuable for decision-making ⚠️ One Line You Can Use in Your Paper

“The observed class imbalance, particularly the low frequency of fatal crashes, reflects real-world crash distributions and highlights the need for careful modeling approaches that account for rare but high-severity outcomes.”

## ***Severity vs Weather***


```
import pandas as pd
import matplotlib.pyplot as plt

# Create contingency table (counts)
weather_severity = pd.crosstab(
    df1['weather_condition_desc'],
    df1['severity_class']
)

# Convert to percentages (row-wise)
weather_severity_pct = weather_severity.div(
    weather_severity.sum(axis=1), axis=0
)
```

```
# Plot
weather_severity_pct.plot(
    kind='bar',
    stacked=True
)

plt.title("Crash Severity Distribution by Weather Condition")
plt.xlabel("Weather Condition")
plt.ylabel("Proportion")
plt.xticks(rotation=45, ha='right')
plt.legend(title="Severity")
plt.show()
```







## Interpretation — Crash Severity by Weather Condition

The distribution of crash severity varies across weather conditions, with notable differences in the proportion of injury-related crashes.



### Key Findings

1. Adverse Weather Increases Injury Proportion Conditions such as sleet/hail, rain, and blowing snow/dirt show: A higher proportion of injury crashes Compared to no adverse conditions



### Interpretation:

Reduced visibility and traction increase the likelihood of injury outcomes.

2. “No Adverse Conditions” Still Has Severe Outcomes Even under normal conditions: Injury and fatal crashes still occur Though property damage dominates



### Insight:

Environmental safety alone does not eliminate crash severity risk.

3. Extreme Conditions Show Volatility Categories like: severe crosswinds sleet hail fog



Show noticeable variation in severity proportions (but likely based on smaller sample sizes)

4. Fatal Crashes Are Rare Across All Conditions Fatal proportion remains consistently low No weather category shows dominance in fatal crashes



### Important:

Fatal events are rare but not strongly concentrated in a single weather condition



### Implication for Your Project

This supports your framework:

Weather is a significant contextual risk factor, particularly for injury-level severity, but does not independently determine fatal outcomes.


## Severity vs Light Condition


```
# Create contingency table
light_severity = pd.crosstab(
    df1['light_condition_desc'],
    df1['severity_class']
)

# Convert to proportions
light_severity_pct = light_severity.div(
    light_severity.sum(axis=1), axis=0
)

# Plot
light_severity_pct.plot(
    kind='bar',
    stacked=True
)

plt.title("Crash Severity Distribution by Light Condition")
plt.xlabel("Light Condition")
plt.ylabel("Proportion")
plt.xticks(rotation=45, ha='right')
plt.legend(title="Severity")
plt.show()
```







## Interpretation — Crash Severity by Light Condition

Crash severity varies notably across lighting conditions, with darker environments associated with a higher proportion of injury-related crashes.



### Key Findings

1. Dark Conditions Increase Injury Severity Both: dark – not lighted dark – lighted show higher injury proportions compared to daylight.



### Interpretation:

Reduced visibility increases crash severity risk, even when artificial lighting is present.

2. Lighting Mitigates but Does Not Eliminate Risk Comparing: dark - lighted vs dark - not lighted



### Difference exists, but:

Injury levels remain elevated in both



### Insight:

Street lighting reduces risk, but does not fully offset the dangers of nighttime driving.

3. Daylight Conditions Are Safer Higher proportion of: Property Damage Only crashes Lower proportion of injury crashes



### Interpretation:

Better visibility reduces severity outcomes

4. Dawn/Dusk Show Transitional Risk Slight variation in injury proportions Not as safe as daylight, not as risky as full darkness

5. Fatal Crashes Remain Rare Across All Conditions No strong concentration in any lighting category 🎯 Implication for Your Project

This is a high-impact variable.

Light condition is one of the strongest predictors of crash severity due to its direct relationship with visibility and driver reaction time.

## ***Severity vs Road Surface Condition***

```


# Create contingency table
road_severity = pd.crosstab(
    df1['road_surface_condition_desc'],
    df1['severity_class']
)

# Convert to proportions
road_severity_pct = road_severity.div(
    road_severity.sum(axis=1), axis=0
)

# Plot
road_severity_pct.plot(
    kind='bar',
    stacked=True
)

plt.title("Crash Severity Distribution by Road Surface Condition")
plt.xlabel("Road Surface Condition")
plt.ylabel("Proportion")
plt.xticks(rotation=45, ha='right')
plt.legend(title="Severity")
plt.show()

```







## Interpretation — Crash Severity by Road Surface Condition

Crash severity varies across road surface conditions, with certain non-dry conditions associated with higher proportions of injury-level crashes.



### Key Findings

1. Slippery Conditions Increase Injury Severity Conditions such as: slush wet water (standing/moving)

show higher proportions of injury crashes compared to dry roads.



### Interpretation:

Reduced traction and vehicle control increase the likelihood of injury outcomes.

2. Dry Conditions Are Relatively Safer Dominated by: Property Damage Only crashes Lower injury proportions compared to adverse surfaces



### Insight:

Stable road conditions reduce crash severity

3. Snow and Ice Show Mixed Effects Snow → relatively lower injury proportion than expected Ice → slightly elevated but not extreme



### Important:

These may reflect:

Driver caution Reduced speeds in severe weather 4. Rare Categories Show Volatility sand/mud, water, etc. High variation due to low sample sizes



### Should be interpreted cautiously

5. Fatal Crashes Remain Rare Across All Conditions No strong concentration in any specific surface type

# EDA SUMMARY

Across three analyses:

-  Weather Moderate effect → impacts injury risk
-  Light Condition Strongest effect → visibility-driven severity
-  Road Surface Moderate effect → traction-driven severity

## ✓ Feature Structuring for Modeling + Web App

This is where we:

- Simplify categories
- Make inputs user-friendly
- Prepare for ML

### Goal

Transform messy categories into:

- Clean
- Interpretable
- UI-friendly (dropdown-ready)

### Key Principle

Your web app user is:

a delivery driver

So inputs must be:

Simple Intuitive Not overly granular

### ***Simplify Weather***

Group logically



```
def simplify_weather(x):
    if x in ['no adverse conditions']:
        return 'clear'
    elif x in ['rain', 'sleet hail']:
        return 'rain'
    elif x in ['snow', 'blowing sand soil dirt or snow']:
        return 'snow'
    elif x in ['fog']:
        return 'fog'
    elif x in ['severe crosswinds']:
        return 'wind'
    else:
        return 'other'

df1['weather_simple'] = df1['weather_condition_desc'].apply(simplify_
```

```
def simplify_light(x):
    if 'daylight' in x:
        return 'day'
    elif 'dark' in x:
        return 'night'
    elif x in ['dawn', 'dusk']:
        return 'twilight'
    else:
        return 'unknown'

df1['light_simple'] = df1['light_condition_desc'].apply(simplify_ligh
```

```
def simplify_road(x):
    if x == 'dry':
        return 'dry'
    elif x in ['wet', 'water(standing moving)']:
        return 'wet'
    elif x in ['snow', 'slush']:
        return 'snow'
    elif x == 'ice':
        return 'ice'
    else:
        return 'other'

df1['road_simple'] = df1['road_surface_condition_desc'].apply(simplif
```

```
def simplify_traffic(x):  
    if 'not divided' in x:  
        return 'undivided'  
    elif 'divided' in x:  
        return 'divided'  
    elif 'one-way' in x:  
        return 'one_way'  
    else:  
        return 'other'  
  
df1['traffic_simple'] = df1['trafficway_desc'].apply(simplify_traffic)
```

```
df1[['weather_simple', 'light_simple', 'road_simple', 'traffic_simple']  
  
for col in ['weather_simple', 'light_simple', 'road_simple', 'traffic_  
    print(f"\n{col}")  
    print(df1[col].value_counts())
```

```
weather_simple  
weather_simple  
clear      1791  
rain       229  
other      161  
snow       99  
fog        21  
wind       10  
Name: count, dtype: int64
```

```
light_simple  
light_simple  
day        1456  
night      763  
twilight   85  
unknown    7  
Name: count, dtype: int64
```

```
road_simple  
road_simple  
dry        1806  
wet        325  
snow       95  
ice        68  
other      17  
Name: count, dtype: int64
```

```
traffic_simple  
traffic_simple  
divided    1402  
undivided  636  
other      273  
Name: count, dtype: int64
```

### *production-ready inputs*

| Feature        | Type        | Web App Ready? |
|----------------|-------------|----------------|
| weather_simple | categorical | ✓              |
| light_simple   | categorical | ✓              |
| road_simple    | categorical | ✓              |
| traffic_simple | categorical | ✓              |

## ✓ Prepare Data for Modeling



Create a clean modeling dataset with:

- Features (X)
- Target (y)

```
# Select features
features = [
    'weather_simple',
    'light_simple',
    'road_simple',
    'traffic_simple'
]

X = df1[features]
y = df1['severity_class']

# Quick check
print(X.head())
print(y.head())
```

```
   weather_simple light_simple road_simple traffic_simple
0         clear          day         dry      undivided
1         clear          day         dry      undivided
2         clear        night         dry       divided
3         clear          day         dry       divided
4         clear          day         dry      undivided
0   Property_Damage
1   Property_Damage
2   Property_Damage
3   Property_Damage
4         Injury
Name: severity_class, dtype: category
Categories (3, object): ['Fatal', 'Injury', 'Property_Damage']
```

```
X_encoded = pd.get_dummies(X, drop_first=True)
```

```
print(X_encoded.shape)
print(X_encoded.head())
```

```
(2311, 14)
  weather_simple_fog  weather_simple_other  weather_simple_rain  \
0                False                False                False
1                False                False                False
2                False                False                False
3                False                False                False
4                False                False                False

  weather_simple_snow  weather_simple_wind  light_simple_night  \
0                False                False                False
1                False                False                False
2                False                False                True
3                False                False                False
4                False                False                False

  light_simple_twilight  light_simple_unknown  road_simple_ice  \
0                False                False                False
1                False                False                False
2                False                False                False
3                False                False                False
4                False                False                False

  road_simple_other  road_simple_snow  road_simple_wet  traffic_simple
0                False                False                False
1                False                False                False
2                False                False                False
3                False                False                False
4                False                False                False

  traffic_simple_undivided
0                True
1                True
2                False
3                False
4                True
```

### 👉 Interpretation:

- 2311 observations ✓
- 14 encoded features ✓
- No explosion in dimensionality ✓
- Encoding Check

You now have:

- Proper one-hot encoding
- No redundant columns (drop\_first=True) ✓

Example:

- weather\_simple\_clear is baseline
- Others are comparisons → correct

### 🧠 What This Means

You now have a fully model-ready dataset:

#### ✓ Features (X\_encoded)

- Clean
- Numerical
- Interpretable

#### ✓ Target (y)

- Multi-class severity

👉 this is the ML pipeline

## ✓ **MODEL BUILDING AND TESTING**

### ***Multiclass Logistic Regression (Baseline Model)***

Add blockquote

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X_encoded, y, test_size=0.2, random_state=42, stratify=y
)

# Model
model = LogisticRegression(max_iter=1000, multi_class='multinomial')

# Train
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)

/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic
warnings.warn(

```

```

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

```

|                 | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Fatal           | 0.00      | 0.00   | 0.00     | 10      |
| Injury          | 0.00      | 0.00   | 0.00     | 167     |
| Property_Damage | 0.62      | 1.00   | 0.76     | 286     |
| accuracy        |           |        | 0.62     | 463     |
| macro avg       | 0.21      | 0.33   | 0.25     | 463     |
| weighted avg    | 0.38      | 0.62   | 0.47     | 463     |

```

[[ 0  0 10]
 [ 0  0 167]
 [ 0  0 286]]

```

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result

```



 Interpretation — Your Model FAILED (and that's GOOD)  What Happened  
 Confusion Matrix: Predicted → ONLY Property\_Damage Classification Report: Fatal →  
~~✗~~ not predicted at all Injury → ~~✗~~ not predicted at all Property Damage → ✓  
 predicted everything  Root Cause (CRITICAL INSIGHT)  Class Imbalance Problem

Your data:

Property Damage → dominant (62%) ~~Injury → moderate (36%) Fatal → tiny (~2%)~~   
 What the Model Did

The model learned:

"If I always predict Property Damage, I get 62% accuracy"

 So it ignored everything else

 Why This Is NOT a Failure (Important for your capstone)

This is actually:

Correct behavior from a naive model on imbalanced data

 VERY STRONG LINE FOR YOUR PAPER

"The baseline multinomial logistic regression model collapsed to predicting the majority class exclusively, highlighting the severe class imbalance present in crash severity data and the inadequacy of standard models without imbalance correction."

 What This Means for Your Project

You just discovered:

 Your CORE PROBLEM:

Rare-event prediction (Fatal crashes)

This is:

Real-world ✓ Research-relevant ✓ High-impact ✓

## FIX THE MODEL (CRITICAL STEP)

We now apply:

Class weighting

```
model = LogisticRegression(
    max_iter=1000,
    multi_class='multinomial',
    class_weight='balanced'
)
```

```
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic
warnings.warn(
```

```
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
```

|                 | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Fatal           | 0.03      | 0.50   | 0.05     | 10      |
| Injury          | 0.35      | 0.40   | 0.37     | 167     |
| Property_Damage | 0.61      | 0.16   | 0.26     | 286     |
| accuracy        |           |        | 0.26     | 463     |
| macro avg       | 0.33      | 0.36   | 0.23     | 463     |
| weighted avg    | 0.50      | 0.26   | 0.30     | 463     |

```
[[ 5  4  1]
 [71 67 29]
 [117 122 47]]
```

### ✅ Step 19 Interpretation — Model FIXED (But Not Good Yet)

You successfully broke the majority-class bias.

 What Changed Before: Predicted ONLY Property Damage ❌ Now: Predicts all 3 classes ✅  Key Results  Fatal (Rare Class) Recall = 0.50 ✅ Precision = 0.03 ❌

👉 Meaning:

Model catches 50% of fatal crashes ✓ But predicts many false positives ✗ 🟡 Injury  
(Middle Class) Precision  $\approx 0.35$  Recall  $\approx 0.40$

👉 Decent baseline performance

🟢 Property Damage (Majority) Recall dropped to 0.16 ✗

👉 Tradeoff:

Model sacrifices accuracy on majority class To learn minority classes ⚠️ Accuracy  
Drop (0.62  $\rightarrow$  0.26)

This is expected and correct.

👉 Why?

Model is no longer “cheating” It is actually trying to learn all classes 🧠 CRITICAL  
INSIGHT (This is your project’s core)

You are facing:

Trade-off between detecting severe crashes and overall accuracy

🎯 VERY STRONG LINE FOR YOUR PAPER

“Applying class weighting improved detection of minority classes, particularly fatal crashes, but introduced a trade-off by reducing overall accuracy and increasing false positives, highlighting the inherent difficulty of multi-class severity prediction under imbalanced conditions.”

📌 What This Means for Your System

Your model now:

✓ Detects severe crashes ! But is not precise yet

👉 This is normal at baseline stage

🚀 Step 20 — Improve Model (SMART FIX)

We now make a small but powerful improvement:

👉 Add time feature

Because:

You already saw: Light condition matters Time will reinforce that signal

```
df1['hour'] = df1['report_date'].dt.hour
```

```
features = [  
    'weather_simple',  
    'light_simple',  
    'road_simple',  
    'traffic_simple',  
    'hour'  
]  
  
X = df1[features]
```

```
X_encoded = pd.get_dummies(X, drop_first=True)
```

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix

# Train-test split (IMPORTANT: re-split after new features)
X_train, X_test, y_train, y_test = train_test_split(
    X_encoded, y, test_size=0.2, random_state=42, stratify=y
)

# Model with class balancing
model = LogisticRegression(
    max_iter=1000,
    multi_class='multinomial',
    class_weight='balanced'
)

# Train
model.fit(X_train, y_train)

# Predict
y_pred = model.predict(X_test)

# Evaluate
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

```

```

/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic
warnings.warn(

```

|                 | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Fatal           | 0.03      | 0.50   | 0.05     | 10      |
| Injury          | 0.35      | 0.40   | 0.37     | 167     |
| Property_Damage | 0.61      | 0.16   | 0.26     | 286     |
| accuracy        |           |        | 0.26     | 463     |
| macro avg       | 0.33      | 0.36   | 0.23     | 463     |
| weighted avg    | 0.50      | 0.26   | 0.30     | 463     |

```

[[ 5  4  1]
 [71 67 29]
 [117 122 47]]

```

## 🔴 Step 20 Result — No Improvement 📊 What Happened

Your output is identical to the previous model.

👉 Meaning:

Adding hour had no measurable impact

🧠 Interpretation (Precise, No Guessing) ❌ Time Feature Failed (in current form)

Why?

1. You used raw hour (0–23) Linear model treats it as numeric But time is cyclical, not linear

Example:

23:00 and 01:00 are close in reality Model thinks they are far apart ❌ 2. Light Condition already captures time effect

You already have:

day night twilight

👉 This is a strong proxy for time

So:

Hour added little new information

🎯 CRITICAL INSIGHT (This is important for your paper)

“Temporal effects were largely captured by light condition, and the addition of raw hour as a feature did not improve model performance, suggesting redundancy and the need for more appropriate temporal representations.”

🔴 NEXT STEP — REAL IMPROVEMENT (No Guesswork)

We now fix the actual problem:

Model choice is too simple for this problem

***Random Forest***

```

from sklearn.ensemble import RandomForestClassifier

# Model
rf_model = RandomForestClassifier(
    n_estimators=200,
    random_state=42,
    class_weight='balanced'
)

# Train
rf_model.fit(X_train, y_train)

# Predict
y_pred_rf = rf_model.predict(X_test)

# Evaluate
from sklearn.metrics import classification_report, confusion_matrix

print(classification_report(y_test, y_pred_rf))
print(confusion_matrix(y_test, y_pred_rf))

```

|                 | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Fatal           | 0.03      | 0.80   | 0.05     | 10      |
| Injury          | 0.43      | 0.14   | 0.22     | 167     |
| Property_Damage | 0.65      | 0.21   | 0.32     | 286     |
| accuracy        |           |        | 0.20     | 463     |
| macro avg       | 0.37      | 0.38   | 0.19     | 463     |
| weighted avg    | 0.56      | 0.20   | 0.27     | 463     |

```

[[ 8  0  2]
 [113 24 30]
 [194 32 60]]

```

 Step 21 Interpretation — Random Forest  What Improved  Fatal (Rare Class)  
Recall = 0.80 ✓ (VERY IMPORTANT) Precision = 0.03 ✗

 Meaning:

Model now catches 8 out of 10 fatal crashes ✓ But still flags many false positives ✗

 Injury (Middle Class) Recall dropped to 0.14 ✗

 This is worse than logistic regression → Model is struggling to separate Injury vs Property Damage

🟢 Property Damage Recall = 0.21 ❌

👉 Model is no longer biased—but also not accurate

🔴 Core Problem (Now Very Clear)

Your model is struggling because:

Your features are too coarse to distinguish severity well

🧠 Why This Is Happening (Important Insight)

Current features:

Weather ✓ Light ✓ Road ✓ Traffic ✓

👉 These explain:

environmental risk

BUT NOT:

crash dynamics impact severity interaction complexity 🎯 CRITICAL CAPSTONE INSIGHT

“Environmental variables alone provide limited predictive power for distinguishing crash severity levels, particularly between injury and non-injury outcomes.”

🔴 NEXT STEP — REAL FIX (Feature Upgrade, NOT Model Tuning)

We now improve features, not models.

🔪 Step 22 — Add Interaction Feature (VERY IMPORTANT)

We create:

Combined risk feature

```
# Combine weather + road into one feature
df1['weather_road_combo'] = df1['weather_simple'] + "_" + df1['road_s:
```



```
features = [  
    'weather_simple',  
    'light_simple',  
    'road_simple',  
    'traffic_simple',  
    'weather_road_combo',  
    'hour'  
]  
  
X = df1[features]  
X_encoded = pd.get_dummies(X, drop_first=True)
```

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X_encoded, y, test_size=0.2, random_state=42, stratify=y  
)
```

```
from sklearn.model_selection import train_test_split  
  
X_train, X_test, y_train, y_test = train_test_split(  
    X_encoded, y, test_size=0.2, random_state=42, stratify=y  
)
```

```

from sklearn.ensemble import RandomForestClassifier

# Model
rf_model = RandomForestClassifier(
    n_estimators=200,
    random_state=42,
    class_weight='balanced'
)

# Train
rf_model.fit(X_train, y_train)

# Predict
y_pred_rf = rf_model.predict(X_test)

# Evaluate
from sklearn.metrics import classification_report, confusion_matrix

print(classification_report(y_test, y_pred_rf))
print(confusion_matrix(y_test, y_pred_rf))

```

|                 | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Fatal           | 0.03      | 0.80   | 0.05     | 10      |
| Injury          | 0.42      | 0.14   | 0.21     | 167     |
| Property_Damage | 0.65      | 0.21   | 0.31     | 286     |
| accuracy        |           |        | 0.20     | 463     |
| macro avg       | 0.36      | 0.38   | 0.19     | 463     |
| weighted avg    | 0.55      | 0.20   | 0.27     | 463     |

```

[[ 8  0  2]
 [113 24 30]
 [194 33 59]]

```

## Step 22 Result — NO CHANGE

Your results are identical.

 That means:

The interaction feature did NOT add signal

 **CRITICAL CONCLUSION (THIS IS YOUR PROJECT'S TRUTH)**  Your data cannot strongly predict severity

Not because:

Model is bad ❌ Code is wrong ❌

But because:

The available features do not contain enough information to separate severity classes well

🎯 WHY THIS IS ACTUALLY A STRONG RESULT

This is not failure.

This is:

A valid, research-level finding

🧠 What your pipeline has proven: ✓ Environmental factors: Influence severity (EDA showed this) ! BUT: Are not sufficient for prediction 🔥 THIS IS YOUR CORE CAPSTONE INSIGHT

“While environmental and roadway conditions are associated with variations in crash severity, their standalone predictive power is limited, indicating that severity outcomes are influenced by additional unobserved factors such as driver behavior, vehicle dynamics, and situational context.”

📊 Model Behavior Summary Class Result Meaning Fatal High recall Model overcompensates Injury Weak Not distinguishable Property Damage Weak Tradeoff problem 🚫 DO NOT KEEP TUNING MODELS

At this point:

Changing models → won't fix problem Adding random features → noise

👉 That leads to bad capstone work

🚀 FINAL STEP — POSITION YOUR PROJECT CORRECTLY

We now shift from:

“build best model”

To:

“build decision-support system with probabilistic insight”

🖋️ Step 23 — Use Probabilities Instead of Hard Prediction

This is KEY for your web app.

```
# Get probabilities
probs = rf_model.predict_proba(X_test)

# Convert to DataFrame
prob_df = pd.DataFrame(probs, columns=rf_model.classes_)

prob_df.head()
```

|          | <b>Fatal</b> | <b>Injury</b> | <b>Property_Damage</b> |  |
|----------|--------------|---------------|------------------------|-----------------------------------------------------------------------------------|
| <b>0</b> | 0.000000     | 0.570593      | 0.429407               |                                                                                   |
| <b>1</b> | 0.373052     | 0.326414      | 0.300534               |                                                                                   |
| <b>2</b> | 0.351271     | 0.346384      | 0.302345               |                                                                                   |
| <b>3</b> | 0.351271     | 0.346384      | 0.302345               |                                                                                   |
| <b>4</b> | 0.419367     | 0.362898      | 0.217734               |                                                                                   |

Next steps:

[Generate code with prob\\_df](#)[New interactive sheet](#)

## ✅ Step 23 Validation — SUCCESS

You now have:

Fatal | Injury | Property\_Damage

👉 As probabilities, not hard labels

🧠 What This Means (THIS IS YOUR PRODUCT)

Example row:

Fatal: 0.37

Injury: 0.33

Property Damage: 0.30

👉 Interpretation:

This scenario has a balanced risk across all severity levels

Another row:

Fatal: 0.00

Injury: 0.58

Property Damage: 0.41

👉 Interpretation:

Much higher likelihood of injury outcome

🔴 THIS IS YOUR WEB APP OUTPUT

Not:

“Prediction = Injury”

But:

“Risk Profile”

## ✓ Format Output for Driver (Readable Risk Report)

```

def format_risk_output_v2(row):
    fatal = round(row['Fatal'] * 100, 1)
    injury = round(row['Injury'] * 100, 1)
    pdmg = round(row['Property_Damage'] * 100, 1)

    # Determine highest risk
    risks = {
        "Fatal": fatal,
        "Injury": injury,
        "Minor Damage": pdmg
    }

    main_risk = max(risks, key=risks.get)
    main_value = risks[main_risk]

    print("🚚 DELIVERY RISK SUMMARY")
    print("=====")

    print(f"👉 Primary Risk: {main_risk} ({main_value}%)\n")

    print("Breakdown:")
    print(f" - Fatal: {fatal}%")
    print(f" - Injury: {injury}%")
    print(f" - Minor Damage: {pdmg}%")

    # Action message
    print("\nDriver Guidance:")

    if main_risk == "Fatal":
        print("🔴 HIGH ALERT: High severity risk. Reduce speed and in
    elif main_risk == "Injury":
        print("🟡 CAUTION: Elevated injury risk. Stay alert and avoid
    else:
        print("🟢 NORMAL: No major risk indicators, but remain attent

```

```
format_risk_output_v2(prob_df.iloc[0])
```



## DELIVERY RISK SUMMARY

```
=====
```



Primary Risk: Injury (57.1%)

### Breakdown:

- Fatal: 0.0%
- Injury: 57.1%
- Minor Damage: 42.9%

### Driver Guidance:



CAUTION: Elevated injury risk. Stay alert and avoid aggressive driv

## ✓ FORMAT INPUT FOR DRIVER

```
!pip install ipywidgets
```

```
Requirement already satisfied: pexpect>4.3 in /usr/local/lib/python3.
Requirement already satisfied: notebook>=4.4.1 in /usr/local/lib/pyth
Requirement already satisfied: parso<0.9.0,>=0.8.6 in /usr/local/lib/
Requirement already satisfied: entrypoints in /usr/local/lib/python3.
Requirement already satisfied: jupyter-core>=4.9.2 in /usr/local/lib/
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/l
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/di
Requirement already satisfied: argon2-cffi in /usr/local/lib/python3.
Requirement already satisfied: nbformat in /usr/local/lib/python3.12/
Requirement already satisfied: nbconvert>=5 in /usr/local/lib/python3
Requirement already satisfied: Send2Trash>=1.8.0 in /usr/local/lib/py
Requirement already satisfied: terminado>=0.8.3 in /usr/local/lib/pyt
Requirement already satisfied: prometheus-client in /usr/local/lib/py
Requirement already satisfied: nbclassic>=0.4.7 in /usr/local/lib/pyt
Requirement already satisfied: ptyprocess>=0.5 in /usr/local/lib/pyth
Requirement already satisfied: wcwidth in /usr/local/lib/python3.12/d
Requirement already satisfied: platformdirs>=2.5 in /usr/local/lib/py
Requirement already satisfied: notebook-shim>=0.2.3 in /usr/local/lib
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/pytho
Requirement already satisfied: bleach!=5.0.0 in /usr/local/lib/python
Requirement already satisfied: defusedxml in /usr/local/lib/python3.1
Requirement already satisfied: jupyterlab-pygments in /usr/local/lib/
Requirement already satisfied: markupsafe>=2.0 in /usr/local/lib/pyth
Requirement already satisfied: mistune<4,>=2.0.3 in /usr/local/lib/py
Requirement already satisfied: nbclient>=0.5.0 in /usr/local/lib/pyth
Requirement already satisfied: pandocfilters>=1.4.1 in /usr/local/lib
Requirement already satisfied: fastjsonschema>=2.15 in /usr/local/lib
```

```

Requirement already satisfied: jsonschema>=2.6 in /usr/local/lib/pyth
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/
Requirement already satisfied: argon2-cffi-bindings in /usr/local/lib
Requirement already satisfied: webencodings in /usr/local/lib/python3
Requirement already satisfied: tinycss2>=1.1.0 in /usr/local/lib/pyth
Requirement already satisfied: attrs>=22.2.0 in /usr/local/lib/python
Requirement already satisfied: jsonschema-specifications>=2023.03.6 i
Requirement already satisfied: referencing>=0.28.4 in /usr/local/lib/
Requirement already satisfied: rpds-py>=0.25.0 in /usr/local/lib/pyth
Requirement already satisfied: jupyter-server<3,>=1.8 in /usr/local/l
Requirement already satisfied: cffi>=1.0.1 in /usr/local/lib/python3.
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python
Requirement already satisfied: typing-extensions>=4.0.0 in /usr/local
Requirement already satisfied: pycparser in /usr/local/lib/python3.12
Requirement already satisfied: anyio>=3.1.0 in /usr/local/lib/python3
Requirement already satisfied: jupyter-events>=0.11.0 in /usr/local/l
Requirement already satisfied: jupyter-server-terminals>=0.4.4 in /us
Requirement already satisfied: websocket-client>=1.7 in /usr/local/li
Requirement already satisfied: idna>=2.8 in /usr/local/lib/python3.12
Requirement already satisfied: python-json-logger>=2.0.4 in /usr/loca
Requirement already satisfied: pyyaml>=5.3 in /usr/local/lib/python3.
Requirement already satisfied: rfc3339-validator in /usr/local/lib/py
Requirement already satisfied: rfc3986-validator>=0.1.1 in /usr/local
Requirement already satisfied: fqdn in /usr/local/lib/python3.12/dist
Requirement already satisfied: isoduration in /usr/local/lib/python3.
Requirement already satisfied: jsonpointer>1.13 in /usr/local/lib/pyt
Requirement already satisfied: rfc3987-syntax>=1.1.0 in /usr/local/li
Requirement already satisfied: uri-template in /usr/local/lib/python3
Requirement already satisfied: webcolors>=24.6.0 in /usr/local/lib/py
Requirement already satisfied: lark>=1.2.2 in /usr/local/lib/python3.
Requirement already satisfied: arrow>=0.15.0 in /usr/local/lib/python

```

```

import ipywidgets as widgets
from IPython.display import display

```

## Create Driver Inputs



```
weather_input = widgets Dropdown(
    options=['clear', 'rain', 'snow', 'fog', 'wind', 'other'],
    description='Weather:'
)

light_input = widgets Dropdown(
    options=['day', 'night', 'twilight'],
    description='Light:'
)

road_input = widgets Dropdown(
    options=['dry', 'wet', 'snow', 'ice', 'other'],
    description='Road:'
)

traffic_input = widgets Dropdown(
    options=['divided', 'undivided', 'other'],
    description='Traffic:'
)
```

### ***Prediction Function***

```
def predict_risk(weather, light, road, traffic):

    # --- ADD DEFAULT HOUR (important!) ---
    hour = 12    # assume midday (or we can make it input later)

    # Create input dataframe
    input_df = pd.DataFrame([
        'weather_simple': weather,
        'light_simple': light,
        'road_simple': road,
        'traffic_simple': traffic,
        'weather_road_combo': weather + "_" + road,
        'hour': hour
    ])

    # Encode
    input_encoded = pd.get_dummies(input_df)

    # 🔥 CRITICAL FIX: Align with training columns
    input_encoded = input_encoded.reindex(columns=X_encoded.columns,

    # Predict
    probs = rf_model.predict_proba(input_encoded)

    prob_df_input = pd.DataFrame(probs, columns=rf_model.classes_)

    # Output
    format_risk_output_v2(prob_df_input.iloc[0])
```

```
button = widgets.Button(description="Assess Risk")

def on_button_click(b):
    print("\n")
    predict_risk(
        weather_input.value,
        light_input.value,
        road_input.value,
        traffic_input.value
    )

button.on_click(on_button_click)
```

## ✓ ***BUTTON INPUTS FOR THE DRIVER***

```
display(weather_input, light_input, road_input, traffic_input, button)
```

Weather:

Light:

Road:

Traffic:

Assess Risk

 DELIVERY RISK SUMMARY

=====

 Primary Risk: Minor Damage (89.6%)

Breakdown:

- Fatal: 3.9%
- Injury: 6.6%
- Minor Damage: 89.6%

Driver Guidance:

 NORMAL: No major risk indicators, but remain attentive.



## UPS Delivery Crash Risk Estimator

This app estimates crash severity risk using the final Random Forest workflow from the modeling notebook. Inputs were simplified into driver-facing categories for weather, light, road surface, trafficway type, and hour of day. Important limitation: this model reflects the uploaded SMS crash dataset and has low overall predictive accuracy, despite stronger recall for fatal cases.

Weather

rain

Light Condition

night

Road Surface

wet

Trafficway Type

undivided

Hour of Day

22

0

23

Primary Predicted Severity

Fatal

Risk Summary and Guidance

Primary predicted outcome: Fatal

Driver guidance: High-alert scenario. Reduce speed, extend following distance, and avoid risky maneuvers.

Predicted Probability Breakdown

| Severity        | Probability (%) |
|-----------------|-----------------|
| Fatal           | 64.81           |
| Property_Damage | 20.03           |
| Injury          | 15.17           |

Share via Link

# UPS Delivery Crash Risk Estimator

This app estimates crash severity risk using the final Random Forest workflow from the modeling notebook. Inputs were simplified into driver-facing categories for weather, light, road surface, trafficway type, and hour of day. Important limitation: this model reflects the uploaded SMS crash dataset and has low overall predictive accuracy, despite stronger recall for fatal cases.

Weather

rain

Light Condition

day

Road Surface

wet

Trafficway Type

divided

Hour of Day

01423

Clear

Submit

Primary Predicted Severity

Injury

Risk Summary and Guidance

Primary predicted outcome: Injury

Driver guidance: Elevated injury risk. Stay alert, avoid aggressive lane changes, and be cautious at conflict points.

Predicted Probability Breakdown

| Severity        | Probability (%) |
|-----------------|-----------------|
| Injury          | 57.6            |
| Property_Damage | 42.4            |
| Fatal           | 0               |

Share via Link

# UPS Delivery Crash Risk Estimator

This app estimates crash severity risk using the final Random Forest workflow from the modeling notebook. Inputs were simplified into driver-facing categories for weather, light, road surface, trafficway type, and hour of day. Important limitation: this model reflects the uploaded SMS crash dataset and has low overall predictive accuracy, despite stronger recall for fatal cases.

Weather

wind

Light Condition

day

Road Surface

dry

Trafficway Type

divided

Hour of Day

01123

Clear

Submit

Primary Predicted Severity

Property\_Damage

Risk Summary and Guidance

Primary predicted outcome: Property\_Damage

Driver guidance: Lower severity than the other categories, but caution is still needed.

Predicted Probability Breakdown

| Severity        | Probability (%) |
|-----------------|-----------------|
| Property_Damage | 89.55           |
| Injury          | 6.56            |
| Fatal           | 3.89            |

Share via Link